

# Practice 14: Using Set Operators to Solve Problems

Jesús Manuel Juárez Pasillas

November 18, 2021

## 1 Introduction

The SET operators the results of two queries resulting in a single result. There are four SET operators which all have the same priority, and each one does something different from the rest.

These operators are used to carry out compound queries, with which to obtain results from two or more queries which are joined with one or more SET operators, where the final result will be what all the operators put back.

## 2 Development

### 2.1 Activity 1:

Read all the choices carefully because there might be more than one correct answer.

1. Which of these set operators will not sort the rows?

- A. INTERSECT
- B. MINUS
- C. UNION
- D. UNION ALL

**The UNION ALL operator does not order the result by default, but requires placing the ORDER BY clause at the end of the statement.**

2. Which of these operators will remove duplicate rows from the final result?

- A. INTERSECT
- B. MINUS
- C. UNION

D. UNION ALL

**Only the UNION ALL operator returns all rows regardless of repeating rows.**

3. If a compound query contains both a MINUS and an INTERSECT operator, which will be applied first?
- A. The INTERSECT, because INTERSECT has higher precedence than MINUS.
  - B. The MINUS, because MINUS has a higher precedence than INTERSECT.
  - C. The precedence is determined by the order in which they are specified.
  - D. It is not possible for a compound query to include both MINUS and INTERSECT.

**The Oracle server evaluates the operators from left to right if there are no parentheses indicating the priority.**

4. There are four rows in the REGIONS table. Consider the following statements and choose how many rows will be returned for each: 0, 4, 8, or 16.
- A. select \* from regions union select \* from regions  
Returns 4 rows because it does not return duplicate rows.
  - B. select \* from regions union all select \* from regions 8  
It returns 8 rows because if it returns duplicate rows.
  - C. select \* from regions minus select \* from regions 0  
You don't get any columns because all the rows from the first query are in the second query.
  - D. select \* from regions intersect select \* from regions 4  
Returns 4 rows because the rows returned by the first query are also returned by the second query.
5. Consider this compound query:
- ```
select empno, hired from emp
union all
select emp_id,hired,fired from ex_emp;
```
- A. Because the columns EMPNO and EMP\_ID have different names
  - B. Because the columns EMP.HIRED and EX\_EMP.HIRED are different data types
  - C. Because there are two columns in the first query and three columns in the second query
  - D. For all the reasons above

E. The query will succeed.

**It gives an error because you have to have the same number of columns in the first query as in the second, in addition to the fact that the data types must also match in order and number.**

6. Which line of this statement will cause it to fail?

A. select ename, hired from current\_staff

B. order by ename

C. minus

D. select ename, hired from current staff

E. where deptno=10

F. order by ename;

**You can only have one ORDER BY at the end of the statement.**

7. Study this statement:

select ename from emp

union all

select ename from ex\_emp;

In what order will the rows be returned?

A. The rows from each table will be grouped and within each group will be sorted on ENAME.

B. The rows from each table will be grouped but not sorted.

C. The rows will not be grouped but will all be sorted on ENAME.

D. The rows will be neither grouped nor sorted.

**The UNION ALL operator does not order the result by default, and it also does not group the results of each query.**

## 2.2 Activity 2:

Propose an answer to the following issues:

- How can you present several tables with similar data as one table?
  - Using the UNION or UNION ALL operator depending on whether you want to obtain repeated rows or not. You only have to consult the similar columns of each table (obtaining the same number of columns in the same order by data type), and with this obtain the result as if it were only a table.

- Are there performance issues with compound queries?
  - Compound queries consist of at least two queries to perform the desired operation and depending on the number of unions required, the number of comparisons that will be made with the results of the queries will be greater or less depending on the union operator.

### 2.3 Activity 3:

This exercise must be performed using HR schema.

- a) In this exercise, you will see the effect of the set operators.
- 1 Connect to your database as user HR.
  - 2 Run a query that consult the regions table (region\_name):
    - **Select region\_name from Regions;**

|   | REGION_NAME            |
|---|------------------------|
| 1 | Europe                 |
| 2 | Americas               |
| 3 | Asia                   |
| 4 | Middle East and Africa |

Figure 1: Regions table.

- 3 Query the Regions table twice, using UNION:
  - **Select region\_name from Regions  
UNION  
Select region\_name from Regions;**

|   | REGION_NAME            |
|---|------------------------|
| 1 | Americas               |
| 2 | Asia                   |
| 3 | Europe                 |
| 4 | Middle East and Africa |

Figure 2: Regions table twice, using UNION.

- 4 This time, use UNION ALL:
  - **Select region\_name from Regions  
UNION ALL  
Select region\_name from Regions;**

|   | REGION_NAME            |
|---|------------------------|
| 1 | Europe                 |
| 2 | Americas               |
| 3 | Asia                   |
| 4 | Middle East and Africa |
| 5 | Europe                 |
| 6 | Americas               |
| 7 | Asia                   |
| 8 | Middle East and Africa |

Figure 3: Regions table twice, using UNION ALL.

5 An intersection will retrieve rows common to two queries:

- **Select region\_name from Regions**  
**INTERSECT**  
**Select region\_name from Regions;**

|   | REGION_NAME            |
|---|------------------------|
| 1 | Americas               |
| 2 | Asia                   |
| 3 | Europe                 |
| 4 | Middle East and Africa |

Figure 4: Regions table twice, using INTERSECT.

6 A MINUS will remove common rows:

- **Select region\_name from Regions**  
**MINUS**  
**Select region\_name from Regions;**

|  | REGION_... |
|--|------------|
|--|------------|

Figure 5: Regions table twice, using MINUS.

b) In this exercise, you will run more complex compound queries.

- 1 Connect to your database as user HR.
- 2 Run a simple query to count the employees in three departments (20,30,40), grouped by them:
  - **Select department\_id, count(employee\_id) count\_emp  
from Employees where department\_id in (20,30,40)  
group by department\_id;**

|   | DEPARTMENT_ID | COUNT_EMP |
|---|---------------|-----------|
| 1 | 20            | 2         |
| 2 | 30            | 6         |
| 3 | 40            | 1         |

Figure 6: Count the employees in three departments (20,30,40).

- 3 Obtain the same result with a compound query:
  - **Select department\_id, count(employee\_id)count\_emp from  
Employees group by department\_id  
INTERSECT  
Select department\_id, count(\*) from Employees where  
department\_id in (20,30,40) group by department\_id;**
- 4 Find out (using compound queries) if any managers manage staff in both departments 20 and 30, and exclude any managers with staff in department 40:
  - **Select manager\_id from Employees  
INTERSECT  
Select manager\_id from Employees where  
department\_id = 20  
INTERSECT  
Select manager\_id from Employees where  
department\_id = 30  
INTERSECT  
Select employee\_id from Employees where  
department\_id != 40;**

|   | MANAGER_ID |
|---|------------|
| 1 | 100        |

Figure 7: Managers staff in both departments 20 and 30

- 5 Use a compound query (3 sentences using two set operator) to report salaries (from employees) subtotaled by department (grouped by department\_id), by manager (grouped by manager\_id), and the overall total. Order the query:

- **Select department\_id,null manager\_id,sum(salary)  
from Employees group by department\_id  
UNION  
Select null,manager\_id,sum(salary) from Employees  
group by manager\_id  
UNION ALL  
Select null,null,sum(salary) from Employees;**

|    | DEPARTMENT_ID | MANAGER_ID | SUM(SALARY) |
|----|---------------|------------|-------------|
| 1  | 10            | (null)     | 4400        |
| 2  | 20            | (null)     | 19000       |
| 3  | 30            | (null)     | 24900       |
| 4  | 40            | (null)     | 6500        |
| 5  | 50            | (null)     | 156400      |
| 6  | 60            | (null)     | 28800       |
| 7  | 70            | (null)     | 10000       |
| 8  | 80            | (null)     | 304500      |
| 9  | 90            | (null)     | 58000       |
| 10 | 100           | (null)     | 51608       |
| 11 | 110           | (null)     | 20308       |
| 12 | (null)        | 100        | 155400      |
| 13 | (null)        | 101        | 44916       |
| 14 | (null)        | 102        | 9000        |
| 15 | (null)        | 103        | 19800       |
| 16 | (null)        | 108        | 39600       |

Figure 8: Salaries subtotaled by department, by manager, and the overall total.

- c) Working in the HR schema, design some queries that will generate reports using the set operators. The reports required are as follows:

- 1 Employees have their current job (identified by JOB\_ID) recorded in their EMPLOYEES row. Jobs they have held previously (but not their current job) are recorded in JOB\_HISTORY. Which employees have never changed jobs? The listing should include the employees' EMPLOYEE\_ID and LAST\_NAME.

- **Select employee\_id, last\_name from Employees  
where employee\_id in  
(Select employee\_id from Employees  
MINUS  
Select employee\_id from Job\_History) order by employee\_id;**

|    | EMPLOYEE_ID | LAST_NAME |
|----|-------------|-----------|
| 1  | 100         | King      |
| 2  | 103         | Hunold    |
| 3  | 104         | Ernst     |
| 4  | 105         | Austin    |
| 5  | 106         | Pataballa |
| 6  | 107         | Lorentz   |
| 7  | 108         | Greenberg |
| 8  | 109         | Faviet    |
| 9  | 110         | Chen      |
| 10 | 111         | Sciarra   |
| 11 | 112         | Urman     |
| 12 | 113         | Popp      |
| 13 | 115         | Khoo      |
| 14 | 116         | Baida     |

Figure 9: Employees who have never changed jobs

2 Which employees were recruited into one job, then changed to a different job, but are now back in a job they held before? Again, you will need to construct a query that compares EMPLOYEES with JOB\_HISTORY. The report should show the employees' names and the job titles. Job titles are stored in the table JOBS.

- **Select last\_name, job\_title from Employees  
natural join Jobs natural join  
(Select employee\_id,job\_id from Job\_History  
INTERSECT  
Select employee\_id,job\_id from Employees);**

|   | LAST_NAME | JOB_TITLE                |
|---|-----------|--------------------------|
| 1 | Taylor    | Sales Representative     |
| 2 | Whalen    | Administration Assistant |

Figure 10: Which employees back in a job they held before

3 What jobs has any one employee held? This will be the JOB\_ID for the employee's current job (in EMPLOYEES) and all previous jobs (in JOB\_HISTORY). If the employee has held a job more than once, there is no need to list it more than once. Use a replacement variable to prompt for the EMPLOYEE\_ID and display the job title(s). Employees 101 and 200 will be suitable employees for testing.



- Select job\_title from Jobs natural join  
(Select job\_id from Employees where  
employee\_id = &employee\_id  
UNION  
Select job\_id from Job\_History where  
employee\_id = &employee\_id)

|   | JOB_TITLE                     |
|---|-------------------------------|
| 1 | Public Accountant             |
| 2 | Accounting Manager            |
| 3 | Administration Vice President |

Figure 11: Employee's current job and all previous jobs

## 2.4 Activity 4:

In this activity you will write several queries using the set operators.

1. The HR department needs a list of department IDs for departments that do not contain the job ID ST\_CLERK.
  - Select department\_id from Employees where  
department\_id is not null  
MINUS  
(Select department\_id from Employees  
where job\_id = 'ST\_CLERK'  
UNION  
Select department\_id from Job\_History  
where job\_id = 'ST\_CLERK');

|    | DEPARTMENT_ID |
|----|---------------|
| 1  | 10            |
| 2  | 20            |
| 3  | 30            |
| 4  | 40            |
| 5  | 60            |
| 6  | 70            |
| 7  | 80            |
| 8  | 90            |
| 9  | 100           |
| 10 | 110           |

Figure 12: Departments that do not contain the job ID ST\_CLERK.

2. The HR department needs a list of countries that have no departments located in them. Display the country ID and the name of the countries.

- **Select country\_id, country\_name from Countries  
MINUS  
Select country\_id, country\_name from Countries  
join Locations using (country\_id) join departments  
using (location\_id);**

|    | COUNTRY_ID | COUNTRY_NAME |
|----|------------|--------------|
| 1  | AR         | Argentina    |
| 2  | AU         | Australia    |
| 3  | BE         | Belgium      |
| 4  | BR         | Brazil       |
| 5  | CH         | Switzerland  |
| 6  | CN         | China        |
| 7  | DK         | Denmark      |
| 8  | EG         | Egypt        |
| 9  | FR         | France       |
| 10 | IL         | Israel       |
| 11 | IN         | India        |
| 12 | IT         | Italy        |
| 13 | JP         | Japan        |

Figure 13: Countries that have no departments located in them.

3. Produce a list of jobs for departments 10, 50, and 20, in that order. Display the job ID and department ID by using the set operators.

- **Select distinct job\_id, department\_id from  
Employees where (department\_id=10)  
UNION ALL  
Select distinct job\_id, department\_id from  
Employees where (department\_id=50)  
UNION ALL  
Select distinct job\_id, department\_id from  
Employees where (department\_id=20);**

|   | JOB_ID   | DEPARTMENT_ID |
|---|----------|---------------|
| 1 | AD_ASST  | 10            |
| 2 | ST_MAN   | 50            |
| 3 | SH_CLERK | 50            |
| 4 | ST_CLERK | 50            |
| 5 | MK_MAN   | 20            |
| 6 | MK_REP   | 20            |

Figure 14: Jobs for departments 10, 50, and 20.

4. Create a report that lists the employee IDs and job IDs of those employees who currently have a job title that is the same as their job title when they were initially hired by the company (that is, they changed jobs but have now gone back to doing their original job).

- **Select employee\_id, job\_id from Employees  
INTERSECT  
Select employee\_id, job\_id from Job\_History;**

|   | EMPLOYEE_ID | JOB_ID  |
|---|-------------|---------|
| 1 | 176         | SA_REP  |
| 2 | 200         | AD_ASST |

Figure 15: Employees they changed jobs but have gone back to their original job.

5. The HR department needs a report with the following specifications:
  - Last name and department ID of all employees from the EMPLOYEES table, regardless of whether or not they belong to a department
  - Department ID and department name of all departments from the DEPARTMENTS table, regardless of whether or not they have employees working in them
  - **Select last\_name,department\_id,to\_char(null)  
Department\_name from Employees  
UNION  
Select to\_char(null),department\_id,department\_name  
from Departments;**

| LAST_NAME    | DEPARTMENT_ID | DEPARTMENT_NAME |
|--------------|---------------|-----------------|
| 1 Abel       | 80 (null)     |                 |
| 2 Ande       | 80 (null)     |                 |
| 3 Atkinson   | 50 (null)     |                 |
| 4 Austin     | 60 (null)     |                 |
| 5 Baer       | 70 (null)     |                 |
| 6 Baida      | 30 (null)     |                 |
| 7 Banda      | 80 (null)     |                 |
| 8 Bates      | 80 (null)     |                 |
| 9 Bell       | 50 (null)     |                 |
| 10 Bernstein | 80 (null)     |                 |
| 11 Bissot    | 50 (null)     |                 |

Figure 16: EMPLOYEES.

| LAST_NAME   | DEPARTMENT_ID       | DEPARTMENT_NAME |
|-------------|---------------------|-----------------|
| 104 Weiss   | 50 (null)           |                 |
| 105 Whalen  | 10 (null)           |                 |
| 106 Zlotkey | 80 (null)           |                 |
| 107 (null)  | 10 Administration   |                 |
| 108 (null)  | 20 Marketing        |                 |
| 109 (null)  | 30 Purchasing       |                 |
| 110 (null)  | 40 Human Resources  |                 |
| 111 (null)  | 50 Shipping         |                 |
| 112 (null)  | 60 IT               |                 |
| 113 (null)  | 70 Public Relations |                 |
| 114 (null)  | 80 Sales            |                 |
| 115 (null)  | 90 Executive        |                 |
| 116 (null)  | 100 Finance         |                 |
| 117 (null)  | 110 Accounting      |                 |

Figure 17: DEPARTMENTS.

### 3 Pre-Assessment:

- Practices pre-Assessment for Database Systems Laboratory II

| Practice                                              | Pre-Assessment |
|-------------------------------------------------------|----------------|
| COMPLIES WITH THE REQUESTED FUNCTIONALITY             | X              |
| HAS THE CORRECT INDENTATION                           | X              |
| HAS AN EASY WAY TO ACCESS THE PROVIDED FILES          | X              |
| HAS A REPORT WITH IDC FORMAT                          | X              |
| REPORT INFORMATION IS FREE OF SPELLING ERRORS         | X              |
| DELIVERED IN TIME AND FORM                            | X              |
| IS FULLY COMPLETED (SPECIFY THE PERCENTAGE COMPLETED) | 100%           |

Table 1: Pre-Assessment.

### 4 Conclusion:

The SET operators are very important when trying to compare the results of two or more queries, in which each operator will tell us something different, therefore, it is necessary to use the one that best suits depending on what is required at the time.

It is very important to control the execution of each of these operators since, although they all have the same priority, the sentence will be executed from left to right if there are no parentheses that specify the order.